ROHAN GODHA

rohangodha.com | rgodha3@gatech.edu
/in/rohan-godha | github.com/rgodha24

Experience

Cofounder & CTO, *Tesser (YC W27)* *Nov 2026 - Present*
- Cofounder & CTO of Tesser in Y Combinator's W27 batch

Software Engineer, *Mintlify* *May 2026 - Oct 2026*
- Led redesign & rearchitecture of the Mintlify Editor, including 100x p99 speedups, performance improvements to support 250k+ pages, & caching/localfirst architecture for instant navigation
- Closed the 3 largest deals in Mintlify history through UX feedback, support, and closing incidents

Software Engineering Intern, *Rocketlab* *May 2025 - Aug 2025*
- Replaced manual solar panel validation process with Python based simulation, achieving large speedup through 16-way parallel processing on Kubernetes with RabbitMQ/KEDA auto-scaling
- Improved loadtime of critical factory visualizations by 192x through optimizing PostgreSQL queries

Founder (acquired), *AI UGC Startup* *Feb 2025 - Dec 2025*
- Developed distributed video generation platform using React, TypeScript and AWS Lambda, enabling users to create viral short-form content with AI Voices
- Proved product-market fit by growing my @aws_peter account to 70k+ followers and 4.7M+ views across Instagram, TikTok and YouTube, producing technical tutorials on Docker, AWS, MCP and AI Agents
- Grew SaaS to >$1k MRR by enabling 1500+ creators to generate 25M+ views and 200k+ followers

Education

Georgia Institute of Technology *Aug 2024 - May 2027*
Candidate: B.S. in Computer Science with concentrations in Information Networks and System Architecture
- **Courses**: Database Systems, Computer Organization, Systems and Networks, Data Structures & Algorithms, Design & Analysis of Algorithms
- **Leadership**: Engineering Manager, GT Bits of Good: managing 10+ developers on nonprofit websites

Projects

Stockfish Trainer *rgodha.com/stockfish*
- Doubled GPU throughput by porting Mixture of Experts CUDA kernels from LLMs, writing custom sparse forward and backward TileLang kernels, overlapping H2D transfers, and fusing optimizer steps
- Enabled multi-GPU and multi-node training on SLURM clusters using Ray, PyTorch Lightning, and NCCL
- Increased CPU dataloading & input encoding throughput by 10x through lockfree rust rewrite
- Experimented with applying the Bitter lesson of ML to Stockfish & Chess NNUEs with a MoE architecture

AlphaPaint *rgodha.com/alphapaint*
- Implemented a best first minimax agent in Rust and placed 3rd out of 100+ teams in Bytefight 2026
- Built a custom AlphaZero style training system using Rust, custom lockfree datastructures, CUDA Graphs, PyTorch, and CUDA kernels via TileLang to achieve 100x speedup over naive implementation
- Ran selfplay across 300 parallel CPU cores on SLURM cluster to generate data for a genetic tuner

Open Source Contributions *rgodha.com/oss*
- opencode: 6th earliest contributor; authored the TUI theming system (#3234), headless auth for ChatGPT/Codex (#10890), subagent TUI navigation (#5761), with 9+ total merged PRs
- openai/codex: Bugfixes around bundling (#12415), terminal UX, and API key auth (#4903) in multiple PRs
- SST: Contributed Rust Lambda support (#5432) and sst_sdk crate (#5562) for the SST IaC framework

AlphaSnake *rgodha.com/alphasnake*
- Trained a custom neural network using a genetic algorithm across 96 CPU cores on Google Cloud
- Optimized a game engine using techniques like increasing cache locality, bitmasks, and SIMD
- Achieved 5th place by ELO in the Georgia Tech Bytefight 2025 competition, sponsored by Google

Skills

- **Programming Languages**: Typescript, Rust, Go, Nix, Python, HTML/CSS/Javascript
- **Frameworks & Technologies**: Svelte/Sveltekit, React, Next.js, Tanstack Start, Tailwind CSS, AWS, Infrastructure as Code (IaC), Docker, ray, pytorch, CUDA, SLURM